

# Apache Kafka

# for Java Developers

## Consuming Records

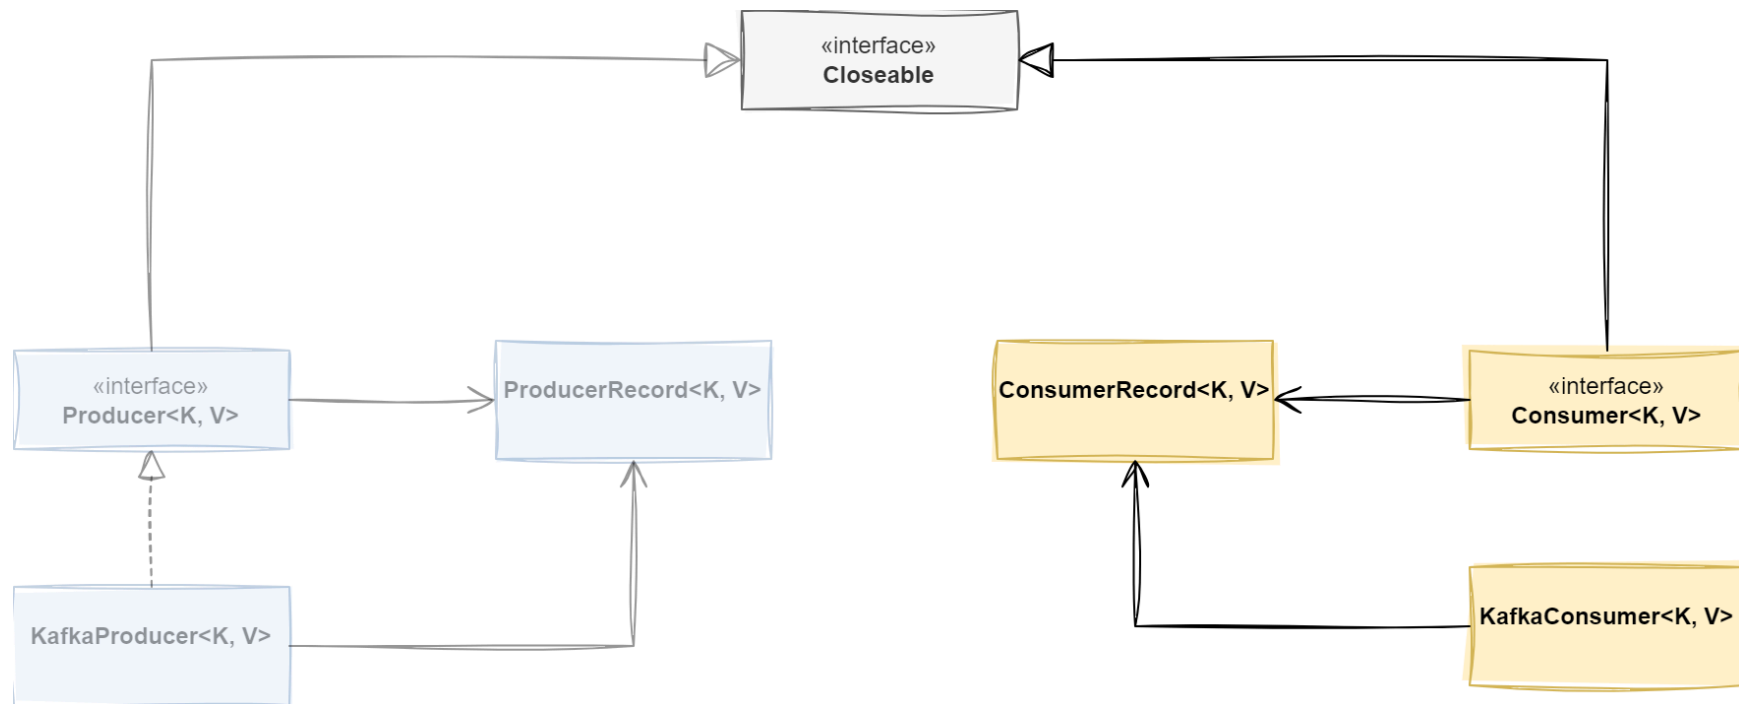
**Boris Fresow, Markus Günther**

JavaLand 2024, Nürburgring

# The Apache Kafka Clients SDK

A Consumer's Perspective

We are able to write a basic consumer by using just these few classes on the right.



## A Consumer<K, V> offers methods for managing subscriptions and consuming records.

```
public interface Consumer<K, V> extends Closeable {  
    // subscription management  
    void subscribe(Collection<String> topics);  
    void subscribe(Collection<String> topics, ConsumerRebalanceListener callback);  
    void subscribe(Pattern pattern);  
    void subscribe(Pattern pattern, ConsumerRebalanceListener callback);  
    void unsubscribe();  
    void assign(Collection<TopicPartition> partitions);  
    Set<TopicPartition> assignment();  
    Set<String> subscription();  
  
    // record consumption  
    ConsumerRecords<K, V> poll(long timeout);  
    ConsumerRecords<K, V> poll(Duration timeout);  
  
    // offset committing (not showing overloaded methods)  
    void commitSync();  
    void commitAsync();  
  
    // manage consumption order (not showing overloaded methods)  
    void seek(TopicPartition partition, long offset);  
    void seekToBeginning(Collection<TopicPartition> partitions);  
    void seekToEnd(Collection<TopicPartition> partitions);  
    // ... and a lot more ...  
}
```

With this knowledge in mind, we are able to write a first, yet simple, consumer!

```
public class BasicConsumer {
    public static void main(String[] args) {
        var topic = "getting-started";
        Map<String, Object> config = Map.of(
            ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092",
            ConsumerConfig.KEY_DESERIALIZER_CONFIG, StringDeserializer.class.getName(),
            ConsumerConfig.VALUE_DESERIALIZER_CONFIG, StringDeserializer.class.getName(),
            ConsumerConfig.GROUP_ID_CONFIG, "basic-consumer-group",
            ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest",
            ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, false);
        try (var consumer = new KafkaConsumer<String, String>(config)) {
            consumer.subscribe(Set.of(topic));
            while (true) {
                var records = consumer.poll(Duration.ofMillis(100));
                for (var record : records) {
                    System.out.println("Received record with value %s\n", record.value());
                }
                consumer.commitAsync();
            }
        }
    }
}
```

# **Subscribing and Consuming**

## A `KafkaConsumer<K, V>` offers two ways for subscribing to the topics.

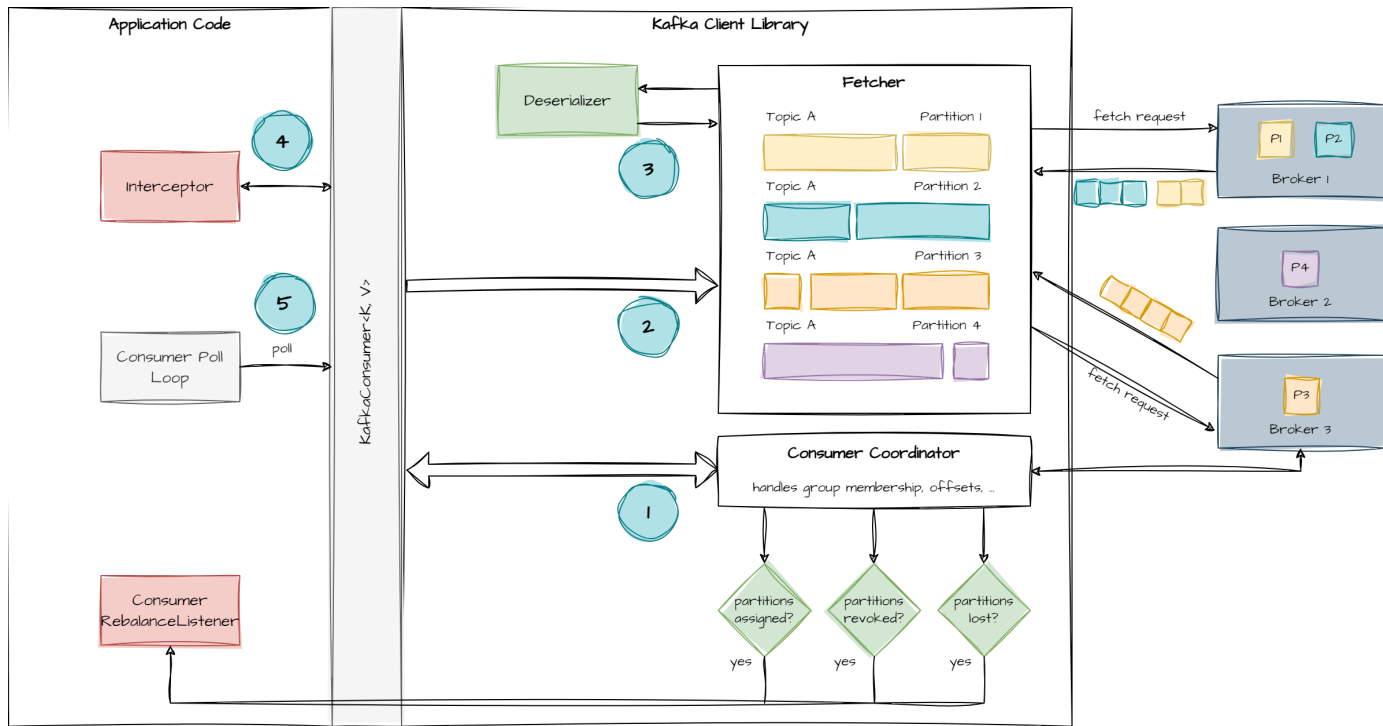
### via `subscribe`

- subscribe to list of topics (regex possible)
- group membership with failure detection
  - client-side
  - server-side
- dynamic partition assignment
- automatic or manual offsets management
- single consumer per partition

### via `assign`

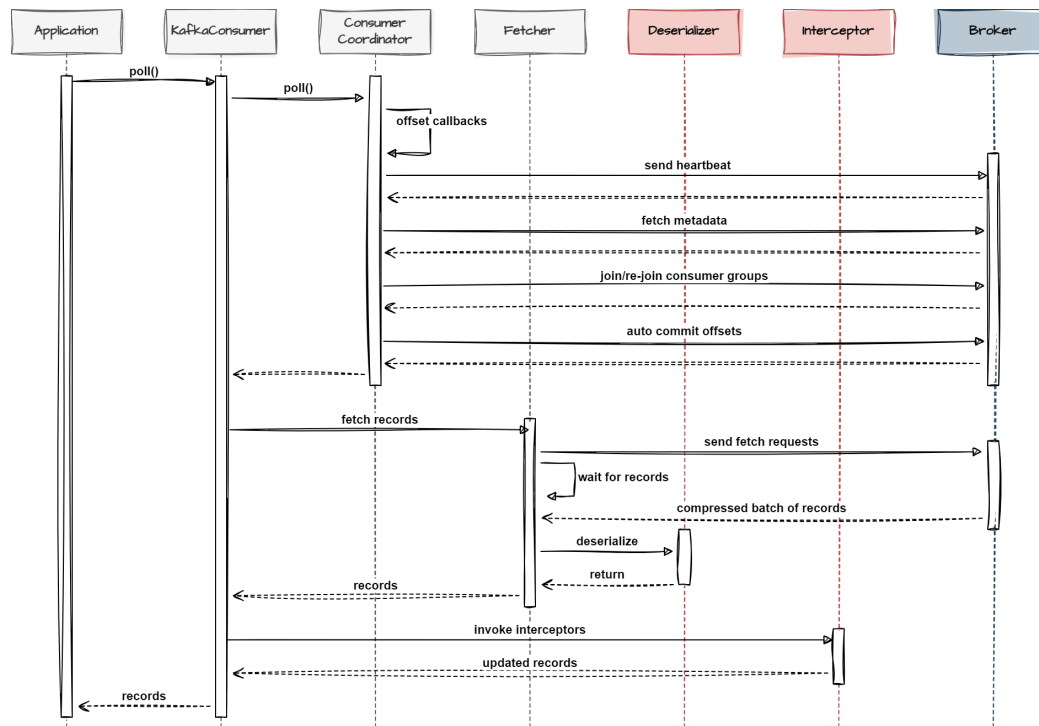
- finer control with topic-partition subscription
- automatic or manual offsets management
- supports multiple consumers per partition

## But what happens after subscribing and entering the poll-loop?

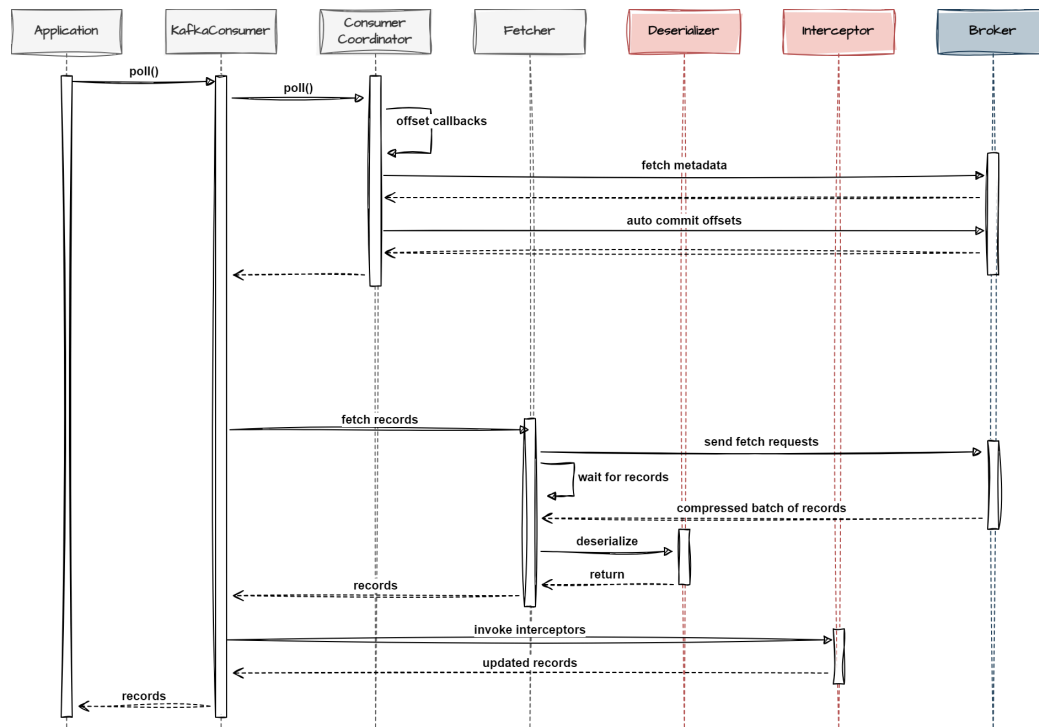




With subscribe, the consumer won't read records until it joins the consumer group.



With assign, the consumer won't invoke group membership functionality.



# Offset Management

**Partition offsets are part of the local consumer state, unless shared with the Kafka cluster.**

- Use consumer offsets as **resumption point**
- Upon re-assigning a consumer, we need to know where to continue
- Skip over any records that have already been processed

## Persisting the consumer state to the Kafka cluster is called committing an offset.

- Commit offset of last record + 1
  1. New consumer joins the consumer group and takes over.
  2. Starts off at the offset of the **last record** that has not been read
  3. This is offset of last record + 1
- Kafka employs a **recursive strategy** when managing offsets
  - Utilizes itself to persist and track offsets
  - cf. topic `__consumer_offsets`

## Controlling when an offset is committed provides flexibility wrt. delivery guarantees.

- Move between *at-most-once* to *at-least-once* simply
  - by committing offsets **before** record processing (at-most-once)
  - by committing offsets **after** record processing completes (at-least-once)

*A committed offset implies that the record **one below that offset** and **all prior records** have been processed by the consumer.*

- Last offset of a batch of records acknowledges the whole batch
- Built your error handling strategy around that fact

By default, a Kafka consumer will automatically commit offsets every five seconds.

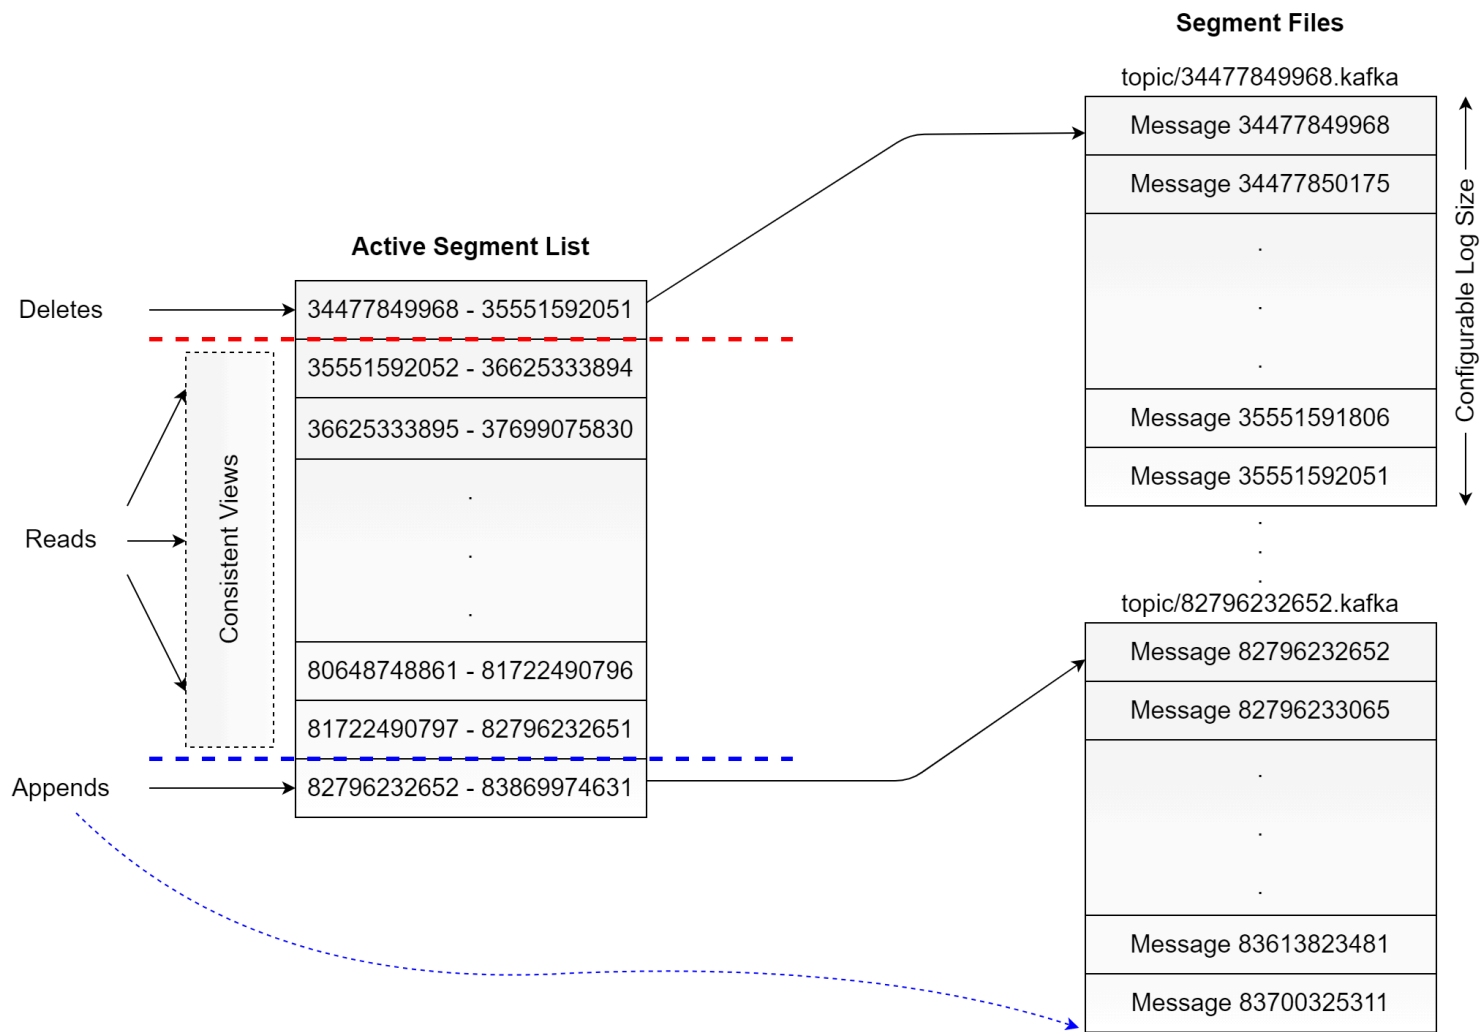
- Adjusting frequency is done by setting `auto.commit.interval.ms`
- Setting `enable.auto.commit` to `false` disables this behavior
- API offers **synchronous** and **asynchronous** commit operations

```
public interface Consumer<K,V> {  
    // only showing commit* methods  
    // synchronous commits  
    void commitSync();  
    void commitSync(Duration timeout);  
    void commitSync(Map<TopicPartition, OffsetAndMetadata> offsets);  
    void commitSync(Map<TopicPartition, OffsetAndMetadata> offsets, Duration timeout);  
    // asynchronous commits  
    void commitAsync();  
    void commitAsync(OffsetCommitCallback callback);  
    void commitAsync(Map<TopicPartition, OffsetAndMetadata> offsets, OffsetCommitCallback callback);  
}
```

In case of no persisted offsets, `auto.offset.reset` controls where the consumer starts.

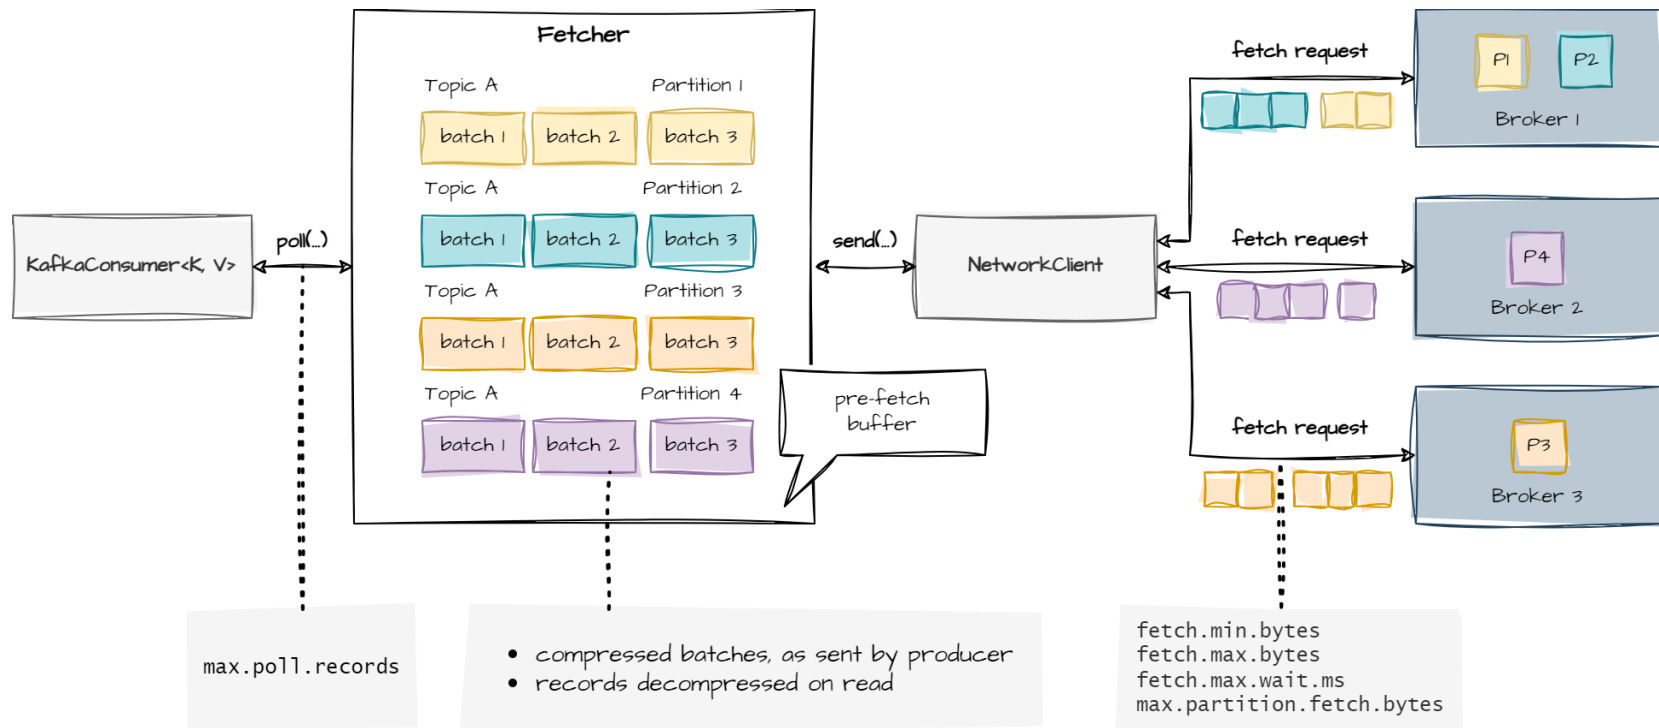
- Offers three different options
  - `earliest`: reset offset to the earliest offset (low watermark)
  - `latest`: reset offset to latest offset (high watermark)
  - `none`: throw exception if there are no offsets present





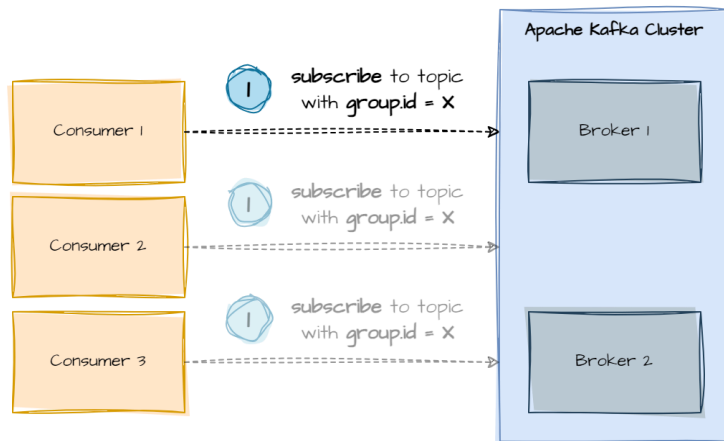
**Fetcher**

# The Kafka consumer uses Fetcher as a buffer that retrieves batches from brokers.

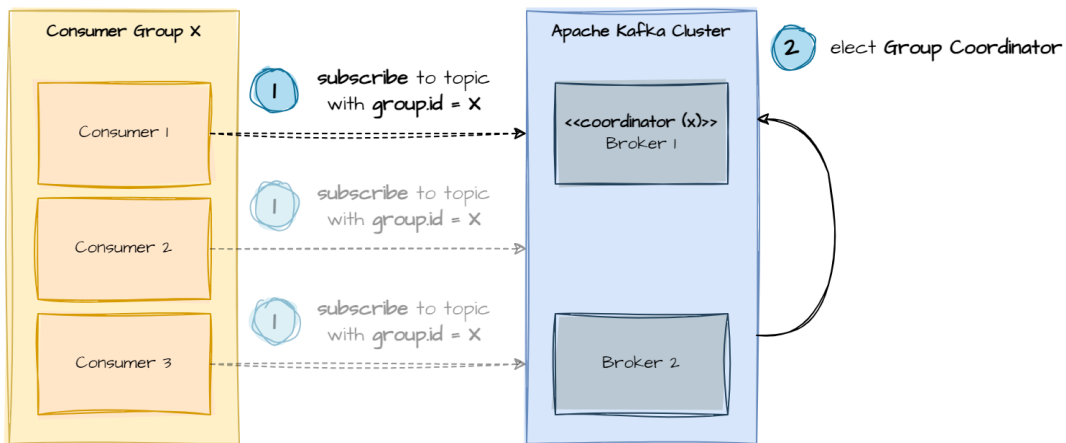


# Consumer Groups

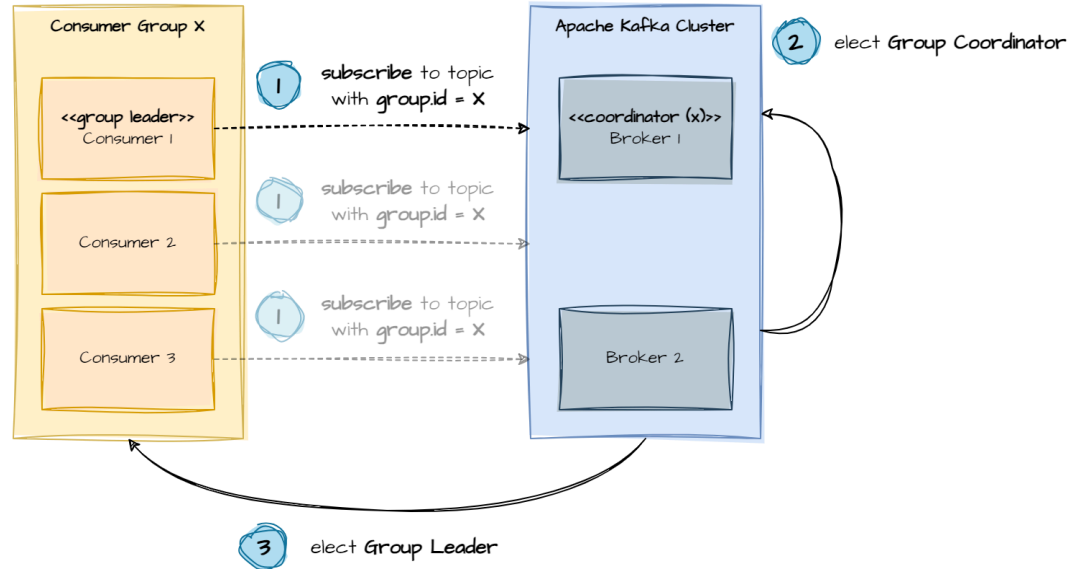
Consumers with the same `group.id` form a consumer group to cooperate.



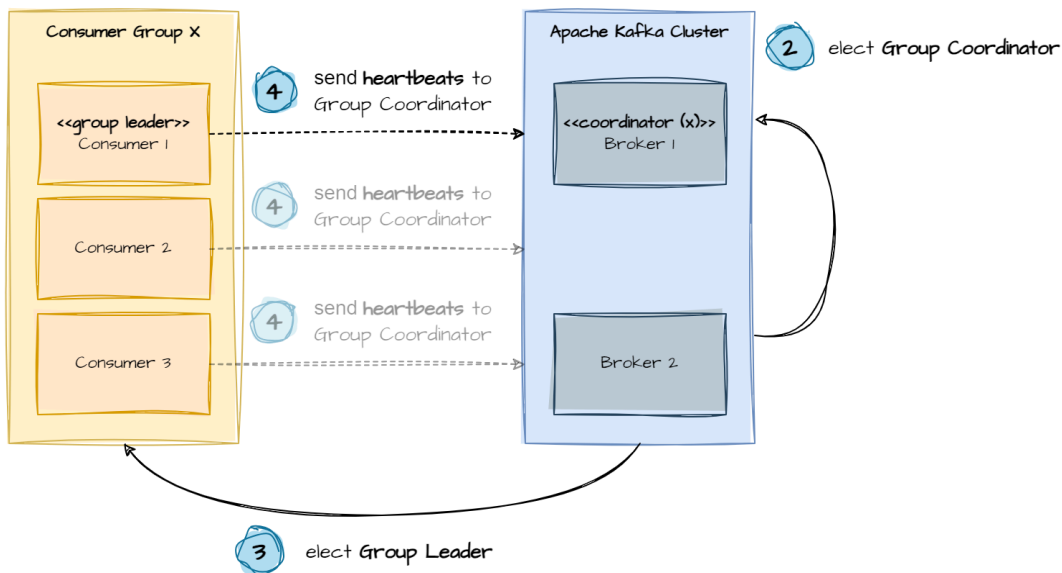
## The Kafka cluster elects one of the brokers as *Group Coordinator*.



## The *Group Coordinator* elects one consumer as *Group Leader*.

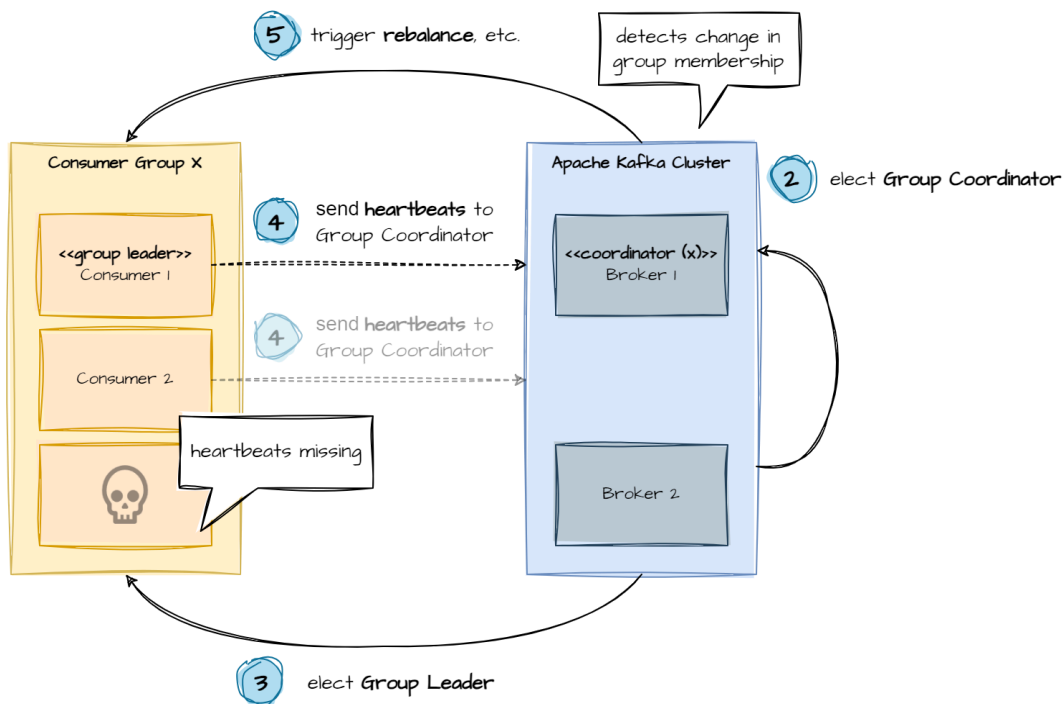


Every consumer of the group sends regular heartbeats to the *Group Coordinator*.





## The *Group Coordinator* detects a change in membership due to missing heartbeats.



**Any change in group membership triggers consumer group rebalances.**

## **Rebalance is triggered when**

- consumer joins the group
- consumer leaves the group
- client-side failure detected via `max.poll.interval.ms`
- server-side failure detected via `session.timeout.ms`

## There are many probable causes for rebalancing.

- Service is scaling up or down
- `poll()` and long message processing occur in the same thread
- Heartbeats do not reach the *Group Coordinator*
- Long JVM garbage collection pauses (stop-the-world)
- Kubernetes pods become CPU-throttled
- Pod evictions due to Kubernetes cluster upgrades
- Networking issues (latency, packet drop, ...)
- ...

The *Group Leader* uses a configurable partition assignment strategy.

- **Range (default)**
  - stop-the-world strategy
  - works on a per-topic basis
  - can generate imbalanced assignments
- **Round Robin**
  - stop-the-world strategy
  - uniformly distributes partitions

## The *Group Leader* uses a configurable partition assignment strategy.

- **Range** (default)
  - stop-the-world strategy
  - works on a per-topic basis
  - can generate imbalanced assignments
- **Round Robin**
  - stop-the-world strategy
  - uniformly distributes partitions
- **Sticky**
  - stop-the-world strategy
  - initial distribution close to *Round Robin*
  - tries to minimize effect of a rebalance
  - can generate imbalanced assignments
- **Cooperative Sticky** (prefer for newer clusters)
  - incremental rebalance
  - does not stop consumption
  - same logic as *Sticky*

## A ConsumerRebalanceListener enables us to react on altered partition assignments.

```
public interface ConsumerRebalanceListener {  
  
    void onPartitionsRevoked(Collection<TopicPartition> partitions);  
  
    void onPartitionsAssigned(Collection<TopicPartition> partitions);  
  
    default onPartitionsLost(Collection<TopicPartition> partitions) {  
        onPartitionsRevoked(partitions);  
    }  
}
```

- Use it to save / restore offsets to / from external storage
- Use it to make sure that outstanding offsets are committed
- Instance is passed to subscribe when subscribing to a topic
- **Important: Different semantics for eager and incremental assignors!**

## What did we learn?

- Client SDK Essentials
- Subscriptions vs. Assignments
- Offset Management
- The `poll()` loop
- Pre-Fetch Buffer
- Consumer Group
- Partition Assignment & Re-Balancing

## What did we learn?

- Client SDK Essentials
- Subscriptions vs. Assignments
- Offset Management
- The `poll()` loop
- Pre-Fetch Buffer
- Consumer Group
- Partition Assignment & Re-Balancing

## What's to follow?

- *Deserialization*
- Interceptors
- Consumer Coordinator
- Consumer Designs
- Transactions



**Questions?**

assignment is available at

[bit.ly/kafka-workshop-exchanging-records](https://bit.ly/kafka-workshop-exchanging-records)